

# Advanced RNGs

---

WISM454 Laboratory Class Scientific Computing, Jan-Willem Buurlage

February 15, 2019

**CWI**

**RNG**


---

# What makes a good RNG?

- Theoretical
  - Period
  - Uniformity
- Empirical
  - Tasks using RNGs, check if output as expected
  - Predictability
  - Invertibility (security!)
- Computational efficiency
- Space and time complexity
- Application specific requirements

# Testing your RNG library

- Basic theory can tell you about e.g. the period, uniformity of your RNG without running it. Also:  **$d$  – dimensional equidistribution**.
- E.g., if LCRNG is used to generate points in  $d$  dimensional space, confined to maximum of  $\sqrt[d]{d!m}$  parallel hyperplanes.  
 $(a, c, m) = (57, 53, 2^7 - 1), \rightarrow 16$  lines



## Software packages ('battery of tests')

- There are software packages such as Diehard (1999) or TestU01 (2007)<sup>1</sup> which contain a collection of statistical tests for your RNGs.
- TestU01 is a C library, can be made part of your code!
- E.g. birthday paradox, rank of random binary matrices, play a game of craps, randomly place spheres in a box, ... all follow a known distribution.

<sup>1</sup><http://simul.iro.umontreal.ca/testu01/tu01.html>

Overview of topics today:

1. Xorshift
2. Linear-feedback shift register
3. Mersenne twister

# Xorshift

```
uint32_t xorshift(uint32_t x) {  
    x ^= x << 13;  
    x ^= x >> 17;  
    x ^= x << 5;  
    return x;  
}
```

## Xorshift^ [George Marsaglia (2003)]

- Consider  $M = \{0, \dots, m-1\}$ , assume that our RNG is still of the form (for  $x_i \in M$ ):

$$x_{i+1} = f(x_i),$$

- Generally, we want  $f$  to be **one-to-one**.
- Usually not enough, since  $f = \text{id}$  is one-to-one but certainly not random.
- Additionally require full period! (Although e.g. ' $f = (+1)$ ' then still works)



# Linear transitions

- Represent  $x_i \in M$  as binary vector:

$$\mathbf{x}^{(i)} = \begin{pmatrix} b_0 \\ b_1 \\ \vdots \\ b_{n-1} \end{pmatrix}$$

take  $f$  to be a linear function, i.e. given by **transition matrix**  $T$ :

$$\mathbf{x}^{(i+1)} = T\mathbf{x}^{(i)}.$$

- $T$  represents one-to-one function iff invertible
- Limit to non-null vectors, i.e. full period consists of all  $m - 1$  non-zero integers in  $M$ .

# Exercise

- **Exercise:** Prove that a non-singular  $T$  generates a non-zero sequence of full period for all non-zero seeds, if and only if the order of  $T$  is  $2^n - 1$  (in group of non-singular  $n \times n$  matrices)

## Intermezzo: bitwise operations

- Bitwise operations are very efficient.

```
uint8_t x = 0b00010010; // 18
```

```
uint8_t y = 0b01010010; // 82
```

- Shifts:

```
x >> 2 // ~> 0b00000100
```

```
x << 2 // ~> 0b01001000
```

- Binary bitwise operations

```
x ^ y // XOR ~> 0b01000000
```

```
x | y // OR ~> 0b01010010
```

```
x & y // AND ~> 0b00010010
```

- We have that  $1 \ll k$  is equal to  $2^k$ , and that XOR is addition modulo two (i.e. addition in our vector space  $(F_2)^n$ )

## How to choose $T$

- Left shift  $L$  (i.e.  $Lx \equiv x \ll 1$ ), right shift  $R$  (i.e.  $Rx \equiv x \gg 1$ ):

$$L = \begin{pmatrix} 0 & \dots & \dots & \dots & 0 \\ 1 & \ddots & \ddots & \ddots & \vdots \\ 0 & 1 & \ddots & \ddots & \vdots \\ \vdots & \ddots & \ddots & \ddots & \vdots \\ 0 & \dots & 0 & 1 & 0 \end{pmatrix}, R = \begin{pmatrix} 0 & 1 & 0 & \dots & 0 \\ \vdots & \ddots & 1 & \ddots & \vdots \\ \vdots & \ddots & \ddots & \ddots & 0 \\ \vdots & \ddots & \ddots & \ddots & 1 \\ 0 & \dots & \dots & \dots & 0 \end{pmatrix}$$

- Note that  $L^a x$ ,  $R^a x$  equal to  $x \ll a$  and  $x \gg a$  respectively.
- Clearly  $L$  and  $R$  singular.

# The xorshift transition matrix

- But:  $\text{Id} + L^a$  and  $\text{Id} + R^a$  non-singular! However, sadly

$$T = (\text{Id} + L^a)(\text{Id} + R^b) \quad (1)$$

does not have the right order for any  $a, b$ .

- Instead choose

$$T = (\text{Id} + L^a)(\text{Id} + R^b)(\text{Id} + L^c) \quad (2)$$

- **Exercise:** Verify experimentally that for (1) no  $a, b$  give  $T$  with required period for  $n = 32$
- **Exercise:** Give all triples  $(a, b, c)$  for which (2) has full period

## Standard xorshift

- The standard Xorshift RNG engine has:

$$T = (Id + L^5)(Id + R^{17})(Id + L^{13}).$$

```
uint32_t xorshift(uint32_t x) {  
    x ^= x << 13;  
    x ^= x >> 17;  
    x ^= x << 5;  
    return x;  
}
```

- There are versions with bigger state and/or more elaborate transition functions that outperform this basic version.

# Linear-feedback shift register

- Xorshift is an example of a linear-feedback shift register

Let  $a \in \{0, 1\}$

## Definition

A (binary) sequence generated by a *shift register* is one satisfying an  $n$ -term recursion:

$$a^{(i+n)} = f(a^{(i)}, \dots, a^{(i+n-1)}).$$

If  $f$  is linear, we speak of a linear-feedback shift register.

# Mersenne twister

- Another example of LFSR
- period of  $2^{19937} - 1$ .
- slower, but  $k$ -equidistributed
- very popular

```
#include <random>
```

```
auto rng = std::mt19937(seed);  
std::cout << rng() << "\n";
```

```
// PS:
```

```
class mt19937 {  
    uint32_t operator()() { ... }  
};
```



- From Wikipedia:

*The Mersenne Twister is the default PRNG for the following software systems: Microsoft Excel,[3] GAUSS,[4] GLib,[5] GNU Multiple Precision Arithmetic Library,[6] GNU Octave,[7] GNU Scientific Library,[8] gretl,[9] IDL,[10] Julia,[11] CMU Common Lisp,[12] Embeddable Common Lisp,[13] Steel Bank Common Lisp,[14] Maple,[15] MATLAB,[16] Free Pascal,[17] PHP,[18] Python,[19][20] R,[21] Ruby,[22] SageMath,[23] Scilab,[24] Stata.[25]*

# Expectations for your RNG library

- Required
  - Usable for randomized algorithms
  - LCRNG
  - Distributions: uniform int, uniform double, Gaussian
- Optional but expected
  - Xorshift
  - Tested with TestU01 / Diehard
  - Benchmarks (random numbers / second)
- Extra credits
  - Mersenne twister
  - Full test-suite based on e.g. TestU01
  - Personal statistical test
  - Other (personal?) engines
  - Extra distributions

# CWI

## C++

---

# Polymorphism (I)

```
class rng {  
    public:  
        virtual int next() = 0;  
};
```

```
class lcrng : public rng {  
    ...  
    int next() override {  
        return ...;  
    }  
    ...  
};
```

- Here, rng is the **base** class and lcrng is the **derived** class that **inherits** from the base class.

## Polymorphism (II)

- An **abstract class** is a class with at least one pure virtual function, like `rng` has:

```
virtual int next() = 0;
```

A **non-abstract class** is also called a **concrete class**

- Objects of an abstract type can not be manipulated by-value, because the representation of an `rng` is unknown. They have to be manipulated using references or pointers:

```
void monte_carlo(rng& engine, ...) { ... }
```

- We call `monte_carlo` with a **concrete** RNG engine. At runtime, the function will call the `next` implementation of this **concrete class\_** (e.g. `lcrng`).

```
auto r = lcrng{14239, 5205, (1 << 30) - 1};  
monte_carlo(r, ...)
```

## Polymorphism (III)

- Abstract classes allow us to leave the choice of e.g. RNG engine to the user, and write our code independently of concrete realizations.

```
class uniform_real_distribution {  
    public:  
        uniform_real_distribution(rng& engine)  
            : engine_(engine) {}  
  
        float sample() {  
            return (float)engine.next() /  
                (engine.max() - 1);  
        }  
  
    private:  
        rng& engine_;  
}
```

## Polymorphism (IV)

```
class abstract {  
    public:  
        virtual void f() = 0;  
};  
  
class concrete : public abstract {  
    public:  
        void f() override {}  
};  
  
abstract a; // ERROR  
concrete b; // fine  
  
void f(abstract& a) {  
    a.f(); // fine  
}
```

## Polymorphism (V)

```
class abstract {  
    public:  
        virtual void f() = 0;  
  
    protected:  
        virtual void g() = 0;  
        int x;  
  
    private:  
        void h() {}  
        int y;  
};
```



## Polymorphism (VI)

- Access specifier, e.g. `public`:

```
class derived : public base ...
```

- private members of base are never visible to derived.
- access specifier specifies maximum visibility of inherited members
- E.g. `class derived : protected base` would make `public` and `protected` members of base, `protected` members of derived.
- For purposes other than inheriting, `protected` is like `private`.

```
class object {  
    public:  
        object() { std::cout << "Constructor\n"; }  
        ~object() { std::cout << "Destructor\n"; }  
};  
  
void f() {  
    object o;  
}  
  
f();  
  
auto o = new object;  
delete o;
```

# Heap storage

- User defined types with heap storage

```
class rng_with_big_state : public rng {  
    public:  
        rng_with_big_state() {  
            state_ = new State;  
        }  
        ~rng_with_big_state() {  
            if (state_) { delete state_; }  
        }  
    private:  
        State* state_;  
};
```

- Hidden new and delete, safer user code:

```
void monte_carlo() {  
    auto r = rng_with_big_state(); // or:  
    rng_with_big_state r; // same thing  
}
```

# Polymorphism and RAI

- Derived classes inherit from base classes.
- Abstract classes versus concrete classes.
- Access specifiers
- RAI allows automatic resource management based on scopes
- These are very important concepts, crucial to understanding how to develop quality C++ software. Spend some time familiarizing yourself with these concepts!
- Any questions/comments on polymorphism and RAI?

# Templates

- Up to now we have discussed runtime polymorphism (also **virtual dispatch**). Templates are 'compile time polymorphism' (**static dispatch**).

```
struct distribution_u32 {  
    uint32_t sample(rng& engine);  
};  
struct distribution_i32 {  
    int32_t sample(rng& engine);  
};  
struct distribution_u64 {  
    uint64_t sample(rng& engine);  
};  
struct distribution_f32 {  
    float sample(rng& engine);  
};
```

- ... there must be an easier way

# Enter templates!

```
template <typename T>
struct distribution {
    virtual T sample(rng& engine) = 0;
};

struct normal_distribution : distribution<float> {
    float sample(rng& engine) override {
        return ...;
    }
};
```

## Fancy tricks!

```
template <typename T,  
typename std::enable_if_t<std::is_floating_point_v<T>>>  
struct normal_distribution : distribution<T> {  
    normal_distribution(T mean, T stddev) { ... }  
  
    T sample(rng& engine) override {  
        return ...;  
    }  
};
```

## Other examples

- Functions can be templates too, compile time values also allowed as **template arguments**.

```
template <int D, typename T>
std::vector<std::array<T, D>> generate_points(
    int count, distribution<T>& f) {
    return ...;
}
```

- Many STL types are templates. We will revisit templates when we discuss the standard library next week!



The logo for CWI (Centrum voor Wiskunde en Informatica) is located in the top left corner. It consists of the letters "CWI" in a white, bold, sans-serif font, set against a red parallelogram background that is wider at the top and tapers towards the bottom.

**CWI**

## **Tutorial**

---

## Common things that are unclear

- ODR
- Calling base constructor
- Virtual functions / vtables

## Concrete things to do

- Implement LCRNG, Xorshift engines
- Implement distributions:
  - Uniform
  - Gaussian (with rejection)
  - Something with inversion
- Write function to randomly permute an array
- Statistically test your generators